

# What is the problem? Generative modeling

- ▶ Data:  $D = \{x_i\}_{i=1}^N$ , where  $x \in \mathcal{X}$  is high-dimensional (e.g., images).
- ▶ Assume  $x_i \sim p_{\text{data}}(x)$ .
- ▶ Goal: learn a probabilistic model  $p_{\theta}(x)$  that approximates  $p_{\text{data}}(x)$  and can generate new samples.
- ▶ Classical objective: maximum likelihood estimation (MLE).

## Maximum likelihood objective

$$\theta^* \in \arg \max_{\theta} \sum_{i=1}^N \log p_{\theta}(x_i).$$

Equivalently (up to constants), minimize negative log-likelihood:

$$\min_{\theta} -\frac{1}{N} \sum_{i=1}^N \log p_{\theta}(x_i).$$

For expressive models, directly specifying  $p_{\theta}(x)$  is hard; latent-variable models provide a structured approach.

## Latent-variable generative models: idea

Introduce latent variables  $z \in \mathcal{Z}$  (typically lower-dimensional):

$$z \sim p_{\theta}(z), \quad x \sim p_{\theta}(x | z).$$

Joint distribution:

$$p_{\theta}(x, z) = p_{\theta}(x | z) p_{\theta}(z).$$


Marginal likelihood (the model for data):

$$p_{\theta}(x) = \int p_{\theta}(x | z) p_{\theta}(z) dz.$$

- ▶  $p_{\theta}(z)$ : prior over latent space.
- ▶  $p_{\theta}(x | z)$ : decoder / generative conditional distribution.

## Why latent variables help

- ▶ High-dimensional  $x$  can be generated from lower-dimensional factors  $z$ .
- ▶ Complex, multimodal distributions over  $x$  can arise from simple priors over  $z$ .
- ▶ Latent space is often smoother and more interpretable (interpolation, clustering, traversals).



# Naive Monte Carlo approximation

Since


$$p_{\theta}(x) = \mathbb{E}_{z \sim p_{\theta}(z)} [p_{\theta}(x | z)],$$

a naive estimator is:

$$p_{\theta}(x) \approx \frac{1}{K} \sum_{k=1}^K p_{\theta}(x | z_k), \quad z_k \sim p_{\theta}(z).$$

Problem:

- ▶ If  $p_{\theta}(x | z)$  is sharply peaked, most  $z_k$  contribute almost nothing.
- ▶ Accurate Monte Carlo may require enormous  $K$ .



## Importance sampling viewpoint

Rewrite the marginal likelihood using any proposal  $q(z | x)$ :

$$p_{\theta}(x) = \int p_{\theta}(x | z) p_{\theta}(z) dz = \int p_{\theta}(x | z) \frac{p_{\theta}(z)}{q(z | x)} q(z | x) dz.$$

Hence:

$$p_{\theta}(x) = \mathbb{E}_{z \sim q(z|x)} \left[ p_{\theta}(x | z) \frac{p_{\theta}(z)}{q(z | x)} \right].$$

Monte Carlo estimator:

$$p_{\theta}(x) \approx \frac{1}{K} \sum_{k=1}^K p_{\theta}(x | z_k) \frac{p_{\theta}(z_k)}{q(z_k | x)}, \quad z_k \sim q(z | x).$$

If  $q(z | x)$  concentrates where  $p_{\theta}(x | z)$  is large, estimation improves.

## ELBO: evidence lower bound

Start from:

$$\log p_{\theta}(x) = \log \int p_{\theta}(x, z) dz.$$

Multiply and divide by  $q_{\phi}(z | x)$ :

$$\log p_{\theta}(x) = \log \int q_{\phi}(z | x) \frac{p_{\theta}(x, z)}{q_{\phi}(z | x)} dz = \log \mathbb{E}_{q_{\phi}(z|x)} \left[ \frac{p_{\theta}(x, z)}{q_{\phi}(z | x)} \right].$$

By Jensen's inequality:

$$\log p_{\theta}(x) \geq \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z | x)].$$

Define the ELBO:

$$\mathcal{L}_{\theta, \phi}(x) = \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x, z) - \log q_{\phi}(z | x)].$$

## Standard ELBO decomposition

Using  $p_{\theta}(x, z) = p_{\theta}(x | z)p_{\theta}(z)$ :

$$\begin{aligned}\mathcal{L}_{\theta, \phi}(x) &= \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x | z)] + \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(z)] - \mathbb{E}_{q_{\phi}(z|x)} [\log q_{\phi}(z | x)] \\ &= \mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x | z)] - D_{\text{KL}}(q_{\phi}(z | x) \| p_{\theta}(z)).\end{aligned}$$

Two parts:

- ▶ Reconstruction (data fit):  $\mathbb{E}_{q_{\phi}(z|x)} [\log p_{\theta}(x | z)]$ .
- ▶ Regularization:  $-D_{\text{KL}}(q_{\phi}(z | x) \| p_{\theta}(z))$ .



## ELBO gap: relation to the true log-likelihood

A key identity:

$$\log p_{\theta}(x) = \mathcal{L}_{\theta, \phi}(x) + D_{\text{KL}}(q_{\phi}(z | x) \| p_{\theta}(z | x)).$$

Since KL is nonnegative:

$$\log p_{\theta}(x) \geq \mathcal{L}_{\theta, \phi}(x).$$

Interpretation:

- ▶ Maximizing ELBO increases  $\log p_{\theta}(x)$ .
- ▶ At optimum, if  $q_{\phi}(z | x) = p_{\theta}(z | x)$ , the bound is tight.

## Training objective over a dataset

Maximize the sum of ELBOs:

$$\max_{\theta, \phi} \sum_{i=1}^N \mathcal{L}_{\theta, \phi}(x_i) = \max_{\theta, \phi} \sum_{i=1}^N \left( \mathbb{E}_{q_{\phi}(z|x_i)} [\log p_{\theta}(x_i | z)] - D_{\text{KL}}(q_{\phi}(z | x_i) \| p_{\theta}(z)) \right).$$

Stochastic optimization uses mini-batches and Monte Carlo estimates of expectations.

## Interpretation of ELBO terms (key ideas)

- **Reconstruction term:**

$$\mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x | z)]$$

encourages the decoder to make  $x$  likely given latents sampled from the encoder.

- **KL regularizer:**

$$D_{\text{KL}}(q_{\phi}(z | x) \parallel p_{\theta}(z))$$

encourages the encoded distribution to match the prior, which:

- makes sampling  $z \sim p(z)$  produce plausible  $x$ ,
- prevents arbitrary, disconnected latent codes.

## VAE architecture: encoder + decoder

VAE combines:

- ▶ **Encoder** (inference network):  $q_{\phi}(z \mid x)$ .
- ▶ **Decoder** (generative network):  $p_{\theta}(x \mid z)$ .
- ▶ **Prior**:  $p_{\theta}(z)$  (often fixed as  $\mathcal{N}(0, I)$ ).

This is “amortized” inference: one neural network outputs parameters of  $q_{\phi}(z \mid x)$  for any  $x$ .

## The gradient problem: sampling inside the encoder

We need gradients of

$$\mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x | z)]$$

w.r.t.  $\phi$ .

But  $z$  is sampled from  $q_{\phi}(z | x)$ , which depends on  $\phi$  (have you ever taken gradient by parameter of distribution??).

Naively,  $\nabla_{\phi} \mathbb{E}_{q_{\phi}}[\cdot]$  is nontrivial (sampling breaks standard backprop).

Solution: **reparameterization trick** for continuous latents (notably Gaussians).

## Reparameterization trick (Gaussian case)

If

$$q_{\phi}(z \mid x) = \mathcal{N}(\mu_{\phi}(x), \text{diag}(\sigma_{\phi}^2(x))),$$

sample by:

$$\varepsilon \sim \mathcal{N}(0, I), \quad z = \mu_{\phi}(x) + \sigma_{\phi}(x) \odot \varepsilon.$$

Then

$$\mathbb{E}_{z \sim q_{\phi}(z|x)}[f(z)] = \mathbb{E}_{\varepsilon \sim \mathcal{N}(0, I)}[f(\mu_{\phi}(x) + \sigma_{\phi}(x) \odot \varepsilon)],$$

which is differentiable w.r.t.  $\phi$  through  $\mu_{\phi}, \sigma_{\phi}$ .

## Monte Carlo estimate of ELBO (per sample)

With reparameterization, using  $K$  noise samples  $\varepsilon_k$ :

$$\mathcal{L}_{\theta,\phi}(x) = \mathbb{E}_{\varepsilon \sim \mathcal{N}(0,I)} [\log p_{\theta}(x \mid z(\varepsilon))] - D_{\text{KL}}(q_{\phi}(z \mid x) \parallel p(z)),$$

approximate:

$$\mathcal{L}_{\theta,\phi}(x) \approx \frac{1}{K} \sum_{k=1}^K \log p_{\theta}(x \mid z_k) - D_{\text{KL}}(q_{\phi}(z \mid x) \parallel p(z)), \quad z_k = \mu_{\phi}(x) + \sigma_{\phi}(x) \odot \varepsilon_k.$$

Commonly,  $K = 1$  works well with mini-batches.

## Connection to classical autoencoders

Deterministic autoencoder:

$$z = f_{\phi}(x), \quad \hat{x} = g_{\theta}(z),$$

with reconstruction loss (e.g.,  $\|x - \hat{x}\|^2$ ).

VAE differences:

- ▶ encoder outputs a *distribution*  $q_{\phi}(z | x)$ , not a point,
- ▶ decoder is probabilistic  $p_{\theta}(x | z)$ ,
- ▶ KL term shapes  $z$  to follow a chosen prior, enabling unconditional sampling.



## Typical training loop (mini-batch SGD)

For a mini-batch  $\{x_b\}_{b=1}^B$ :

1. Encoder forward: compute  $\mu_b = \mu_\phi(x_b)$  and  $\log \sigma_b^2 = \log \sigma_\phi^2(x_b)$ .
2. Sample noise:  $\varepsilon_b \sim \mathcal{N}(0, I)$ .
3. Reparameterize:  $z_b = \mu_b + \sigma_b \odot \varepsilon_b$ .
4. Decoder forward: compute parameters of  $p_\theta(x_b | z_b)$ .
5. Compute:

$$\hat{\mathcal{L}} = \frac{1}{B} \sum_{b=1}^B \log p_\theta(x_b | z_b) - \frac{1}{B} \sum_{b=1}^B D_{\text{KL}}(q_\phi(z | x_b) \| p(z)).$$

6. Update  $\theta, \phi$  by ascending  $\hat{\mathcal{L}}$  (or descending  $-\hat{\mathcal{L}}$ ).

## Generation after training


After training, to generate a new sample:

$$z \sim p(z) \quad (\text{often } \mathcal{N}(0, I)), \quad x \sim p_{\theta}(x \mid z).$$

## Latent space properties (key ideas)

Because  $q_\phi(z | x)$  is regularized toward  $p(z)$ :


- ▶ nearby points in latent space often decode to semantically similar outputs,
- ▶ interpolation is meaningful: for  $z(t) = (1 - t)z_0 + tz_1$
- ▶ latent space can reflect class structure (clusters/regions) when data has discrete modes.



# Manifold



(a) Learned Frey Face manifold



(b) Learned MNIST manifold

## Why the KL term matters for sampling

If we trained only reconstruction (like a deterministic autoencoder):

- ▶ encoded  $z$  could occupy an arbitrary set in  $\mathbb{R}^M$ ,
- ▶ sampling  $z \sim \mathcal{N}(0, I)$  would likely land outside that set,
- ▶ decoded samples would be garbage.

KL regularization enforces:

$$q_{\phi}(z \mid x) \approx p(z),$$

so that sampling from  $p(z)$  produces latents similar to those seen during training.

## Conditional generation: motivation

Sometimes we want to generate not from the overall distribution  $p(x)$ , but from a conditional distribution  $p(x | y)$  for some attribute/label  $y$ :

- ▶ class label (digits, categories),
- ▶ style, domain, speaker identity,
- ▶ any conditioning information available at train/test time.

Solution: **Conditional VAE (CVAE)**.

## Conditional VAE: model and inference

Introduce condition  $y$  and define:

$$p_{\theta}(z \mid y), \quad p_{\theta}(x \mid z, y), \quad q_{\phi}(z \mid x, y).$$

The conditional marginal likelihood:

$$p_{\theta}(x \mid y) = \int p_{\theta}(x \mid z, y) p_{\theta}(z \mid y) dz.$$

ELBO for CVAE:

$$\mathcal{L}_{\theta, \phi}(x, y) = \mathbb{E}_{q_{\phi}(z \mid x, y)}[\log p_{\theta}(x \mid z, y)] - D_{\text{KL}}(q_{\phi}(z \mid x, y) \parallel p_{\theta}(z \mid y)).$$

## CVAE: sampling

To generate conditioned on  $y$ :

$$z \sim p_{\theta}(z \mid y), \quad x \sim p_{\theta}(x \mid z, y).$$

Common choices:

- ▶  $p_{\theta}(z \mid y) = \mathcal{N}(0, I)$  for all  $y$ , but feed  $y$  only to the decoder and encoder;
- ▶ or a class-dependent prior  $p_{\theta}(z \mid y) = \mathcal{N}(\mu_y, \Sigma_y)$  (or NN-parameterized).




## Why CVAE can separate classes better than a vanilla VAE

In a **vanilla VAE**, all data classes share one latent prior:  $p(z) \sim \mathcal{N}(0, I)$ .



In a **Conditional VAE (CVAE)**, we condition the prior on the label  $y$  and effectively give each class its own latent space:

$$p(z \mid y) \sim \mathcal{N}(0, I) \quad \text{for each } y \in \{0, \dots, 9\}.$$



## Why CVAE can separate classes better than a vanilla VAE

As a result, the model does not need to pack all digits into one shared  $p(z)$ ; it can learn *class-specific* latent structure, which often yields cleaner class separation.



## ELBO recap and the full picture

For (unconditional) VAE:

$$\log p_{\theta}(x) \geq \underbrace{\mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x | z)]}_{\text{data fit / reconstruction}} - \underbrace{D_{\text{KL}}(q_{\phi}(z | x) \| p(z))}_{\text{latent regularization}}.$$

Training:

$$\max_{\theta, \phi} \sum_{i=1}^N \mathcal{L}_{\theta, \phi}(x_i).$$

After training:

$$z \sim p(z) \Rightarrow x \sim p_{\theta}(x | z).$$